

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import os, glob, random, itertools
import pandas as pd, numpy as np, matplotlib.pyplot as plt, seaborn as
sns
from tqdm.auto import tqdm
import tensorflow as tf
from tensorflow.keras.preprocessing import
image_dataset_from_directory
```

```
# -----
# 1. Paths & constants
# -----
DATASET_DIR = "/content/drive/MyDrive/ChestXRay2017/chest_xray" #
<-- change if needed
IMG_SIZE     = (224, 224)
BATCH_SIZE   = 32
AUTOTUNE     = tf.data.AUTOTUNE
```

```
# -----
# 2. Build a DataFrame of filepaths + clean labels
# -----
def subclass_from_filename(fname: str) -> str:
    """Return 'bacterial', 'viral', or generic 'pneumonia'."""
    lower = fname.lower()
    if "_bacteria" in lower: return "bacterial"
    if "_virus" in lower: return "viral"
    return "pneumonia" # fallback
```

```
records = []
for split in ("train", "test"):
    for cls in ("NORMAL", "PNEUMONIA"):
        folder = os.path.join(DATASET_DIR, split, cls)
        for fp in glob.glob(os.path.join(folder, "*.jpeg")):
            label = "normal" if cls == "NORMAL" else
subclass_from_filename(fp)
            records.append((fp, label, split))
```

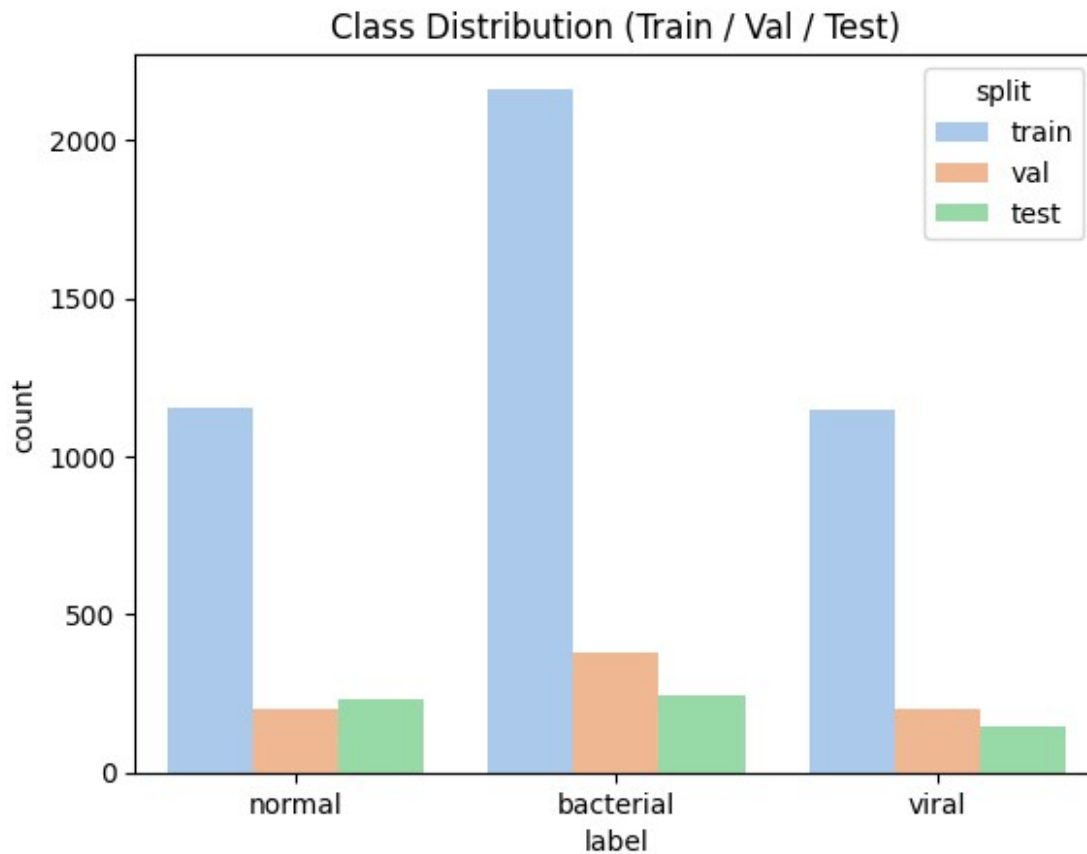
```
df = pd.DataFrame(records, columns=["filepath", "label", "split"])
print("Total images:", len(df))
df.head()
```

Total images: 5876

```
{"summary": "{\n  \"name\": \"df\", \n  \"rows\": 5876, \n  \"fields\": \n[\n  {\n    \"column\": \"filepath\", \n    \"properties\": {\n
```



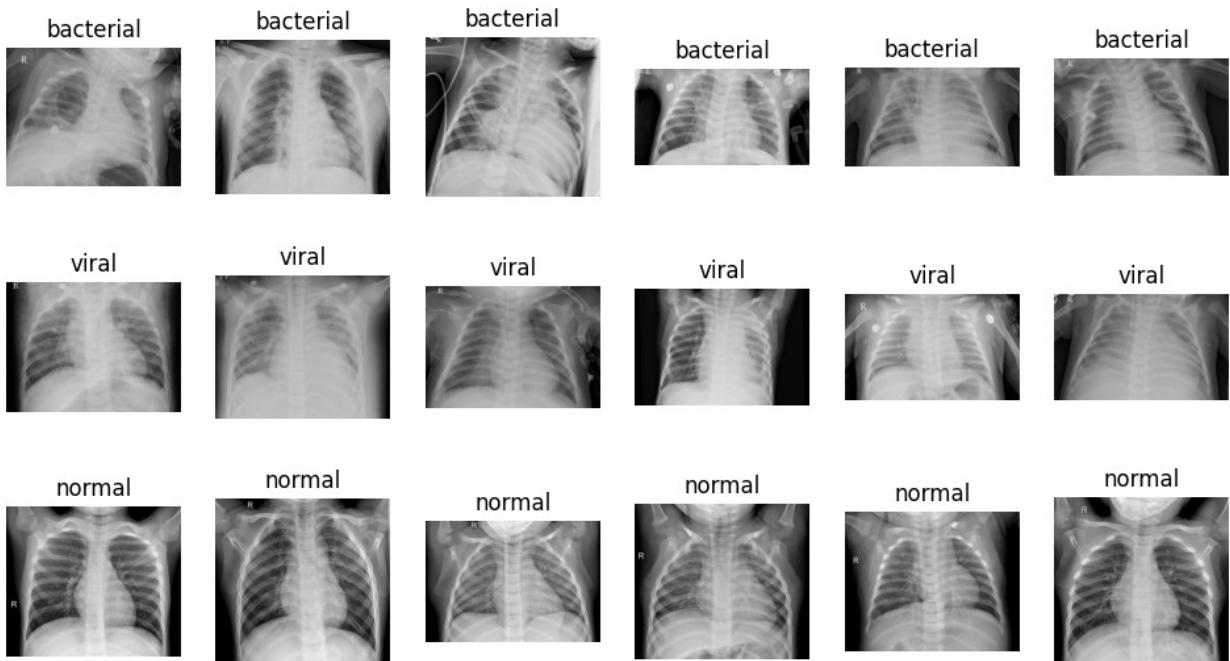
```
plt.title("Class Distribution (Train / Val / Test)")
plt.show()
```



## Random image mosaics for sanity-check

```
# Random image mosaics for sanity-check
def show_random_samples(lbl="bacterial", n=6):
    sample_paths = df[(df.label==lbl) & (df.split=="train")]
    ["filepath"].sample(n)
    plt.figure(figsize=(12,2))
    for i, fp in enumerate(sample_paths, 1):
        img = plt.imread(fp)
        plt.subplot(1, n, i)
        plt.imshow(img, cmap="gray")
        plt.axis("off")
        plt.title(lbl)
    plt.show()

for lbl in ["bacterial", "viral", "normal"]:
    show_random_samples(lbl)
```



```
# Image-size statistics
shapes = [tf.image.decode_jpeg(tf.io.read_file(fp), channels=3).shape
          for fp in tqdm(df["filepath"])]
shape_df = pd.DataFrame(shapes, columns=["h", "w", "c"])
print(shape_df.describe()[["h", "w"]])

{"model_id": "759a0409fe904146811c39ffb004b0e2", "version_major": 2, "version_minor": 0}
```

	h	w
count	5876.000000	5876.000000
mean	970.44418	1327.361811
std	383.36124	363.491567
min	127.000000	384.000000
25%	688.000000	1056.000000
50%	887.000000	1280.000000
75%	1187.000000	1560.000000
max	2713.000000	2916.000000

## Preprocessing

```
# list of filepaths & labels → tf.data.Dataset
label_to_index = {lbl:i for i, lbl in
                   enumerate(sorted(df["label"].unique()))}
```

```

def paths_to_dataset(filepaths, labels):
    ds = tf.data.Dataset.from_tensor_slices((filepaths, labels))
    # -- map strings → image tensors + one-hot labels
    def _load(path, lbl):
        img = tf.io.read_file(path)
        img = tf.image.decode_jpeg(img, channels=3)
        img = tf.image.resize(img, IMG_SIZE)
        img = tf.cast(img, tf.float32) / 255.0 # [0,1]
        lbl = tf.one_hot(lbl, depth=len(label_to_index))
        return img, lbl
    ds = ds.map(_load, num_parallel_calls=AUTOTUNE)
    return ds

# Gather splits
train_paths = df[df.split=="train"]["filepath"].tolist()
train_labels = df[df.split=="train"]
["label"].map(label_to_index).tolist()

val_paths = df[df.split=="val"]["filepath"].tolist()
val_labels = df[df.split=="val"]
["label"].map(label_to_index).tolist()

test_paths = df[df.split=="test"]["filepath"].tolist()
test_labels = df[df.split=="test"]
["label"].map(label_to_index).tolist()

# Build datasets
train_ds = paths_to_dataset(train_paths, train_labels)\
    .shuffle(1024).batch(BATCH_SIZE).prefetch(AUTOTUNE)

val_ds = paths_to_dataset(val_paths, val_labels)\
    .batch(BATCH_SIZE).prefetch(AUTOTUNE)

test_ds = paths_to_dataset(test_paths, test_labels)\
    .batch(BATCH_SIZE).prefetch(AUTOTUNE)

print("Dataset shapes →", next(iter(train_ds))[0].shape,
      next(iter(train_ds))[1].shape)

Dataset shapes → (32, 224, 224, 3) (32, 3)

```

## Model Implementation

```

import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix

```

```

import seaborn as sns
import time

# Dataset-specific values
NUM_CLASSES = 3
EPOCHS = 15
INPUT_SHAPE = (224, 224, 3)
class_names = ['WithMask', 'WithoutMask', 'MaskWearedIncorrect']

```

We Build Four models The four models include:

- DenseNet121 (Transfer Learning)
- Custom CNN (lightweight from scratch)
- MobileNetV2 (Transfer Learning)
- VGG16 (Transfer Learning)

```

# -----
# 1. DenseNet121 Model
# -----
def build_densenet_model(num_classes):
    base = tf.keras.applications.DenseNet121(
        include_top=False, weights="imagenet",
        input_shape=INPUT_SHAPE, pooling="avg")
    base.trainable = False
    inputs = layers.Input(shape=INPUT_SHAPE)
    x = base(inputs, training=False)
    x = layers.Dropout(0.3)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    return models.Model(inputs, outputs)

# -----
# 2. Custom CNN Model
# -----
def build_custom_cnn(num_classes):
    inputs = layers.Input(shape=INPUT_SHAPE)
    x = layers.Conv2D(32, 3, activation='relu', padding='same')(inputs)
    x = layers.MaxPooling2D()(x)
    x = layers.Conv2D(64, 3, activation='relu', padding='same')(x)
    x = layers.MaxPooling2D()(x)
    x = layers.Conv2D(128, 3, activation='relu', padding='same')(x)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dropout(0.3)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    return models.Model(inputs, outputs)

```

```

# -----
# 3. MobileNetV2 Model
# -----
def build_mobilenet_model(num_classes):
    base = tf.keras.applications.MobileNetV2(
        include_top=False, weights="imagenet",
        input_shape=INPUT_SHAPE, pooling="avg")
    base.trainable = False
    inputs = layers.Input(shape=INPUT_SHAPE)
    x = base(inputs, training=False)
    x = layers.Dropout(0.3)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    return models.Model(inputs, outputs)

# -----
# 4. VGG16 Model
# -----
def build_vgg16_model(num_classes):
    base = tf.keras.applications.VGG16(
        include_top=False, weights="imagenet",
        input_shape=INPUT_SHAPE, pooling="avg")
    base.trainable = False
    inputs = layers.Input(shape=INPUT_SHAPE)
    x = base(inputs, training=False)
    x = layers.Dropout(0.3)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    return models.Model(inputs, outputs)

```

Train and Evaluate the model

```

# Compile and train
def compile_and_train(model, name, train_ds, val_ds, epochs=EPOCHS):
    model.compile(
        optimizer=tf.keras.optimizers.Adam(1e-4),
        loss='categorical_crossentropy',
        metrics=['accuracy',
                 tf.keras.metrics.AUC(name='auc'),
                 tf.keras.metrics.Precision(name='precision'),
                 tf.keras.metrics.Recall(name='recall')]
    )
    print(f"\n Training {name}...")
    start_time = time.time()
    history = model.fit(train_ds, validation_data=val_ds,
                        epochs=epochs, verbose=2)
    train_time = time.time() - start_time
    print(f" Training time for {name}: {train_time:.2f} seconds")
    return model, history, train_time

```

```

# Train and evaluate all models
models_to_train = {
    "DenseNet121": build_densenet_model(NUM_CLASSES),
    "Custom CNN": build_custom_cnn(NUM_CLASSES),
    "MobileNetV2": build_mobilenet_model(NUM_CLASSES),
    "VGG16": build_vgg16_model(NUM_CLASSES)
}

model_histories = {}

for model_name, model_instance in models_to_train.items():
    model, history, train_time = compile_and_train(model_instance,
    model_name, train_ds, val_ds)

    # Evaluation
    print(f"\n Evaluating {model_name}...")
    start_test = time.time()
    y_true, y_pred = [], []

    for x_batch, y_batch in test_ds:
        preds = model.predict(x_batch)
        y_pred.extend(np.argmax(preds, axis=1))
        y_true.extend(np.argmax(y_batch.numpy(), axis=1))

    test_time = time.time() - start_test
    print(f" Testing time: {test_time:.2f} seconds")

    print(" Classification Report:")
    print(classification_report(y_true, y_pred,
    target_names=class_names))

    # Confusion Matrix
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
        xticklabels=class_names, yticklabels=class_names)
    plt.title(f"{model_name} - Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.tight_layout()
    plt.show()

    # Training history
    plt.figure(figsize=(12, 4))
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Val Accuracy')
    plt.title(f"{model_name} - Accuracy")
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')

```



```

plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title(f'{model_name} - Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/densenet/>

densenet121\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5  
29084464/29084464 \_\_\_\_\_ 2s 0us/step

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/)

mobilenet\_v2\_weights\_tf\_dim\_ordering\_tf\_kernels\_1.0\_224\_no\_top.h5  
9406464/9406464 \_\_\_\_\_ 2s 0us/step

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

58889256/58889256 \_\_\_\_\_ 4s 0us/step

\n Training DenseNet121...

Epoch 1/15

140/140 - 1547s - 11s/step - accuracy: 0.4301 - auc: 0.6179 - loss: 1.2300 - precision: 0.4524 - recall: 0.3564 - val\_accuracy: 0.4365 - val\_auc: 0.6027 - val\_loss: 1.2767 - val\_precision: 0.4407 - val\_recall: 0.3769

Epoch 2/15

140/140 - 50s - 357ms/step - accuracy: 0.4321 - auc: 0.6144 - loss: 1.2291 - precision: 0.4645 - recall: 0.3560 - val\_accuracy: 0.4721 - val\_auc: 0.6587 - val\_loss: 1.1274 - val\_precision: 0.5030 - val\_recall: 0.4213

Epoch 3/15

140/140 - 45s - 321ms/step - accuracy: 0.4899 - auc: 0.6782 - loss: 1.0881 - precision: 0.5286 - recall: 0.4097 - val\_accuracy: 0.5025 - val\_auc: 0.7006 - val\_loss: 1.0244 - val\_precision: 0.5374 - val\_recall: 0.4467

Epoch 4/15

140/140 - 50s - 360ms/step - accuracy: 0.5296 - auc: 0.7209 - loss: 1.0041 - precision: 0.5793 - recall: 0.4568 - val\_accuracy: 0.5152 - val\_auc: 0.7318 - val\_loss: 0.9478 - val\_precision: 0.5784 - val\_recall: 0.4683

Epoch 5/15

140/140 - 50s - 360ms/step - accuracy: 0.5690 - auc: 0.7535 - loss: 0.9410 - precision: 0.6089 - recall: 0.4908 - val\_accuracy: 0.5355 - val\_auc: 0.7553 - val\_loss: 0.9023 - val\_precision: 0.5959 - val\_recall: 0.4848

Epoch 6/15

140/140 - 45s - 325ms/step - accuracy: 0.5871 - auc: 0.7734 - loss: 0.9007 - precision: 0.6313 - recall: 0.5085 - val\_accuracy: 0.5660 - val\_auc: 0.7864 - val\_loss: 0.8393 - val\_precision: 0.6287 - val\_recall: 0.4898

Epoch 7/15

140/140 - 86s - 616ms/step - accuracy: 0.6120 - auc: 0.7947 - loss: 0.8600 - precision: 0.6507 - recall: 0.5341 - val\_accuracy: 0.6041 - val\_auc: 0.8020 - val\_loss: 0.8104 - val\_precision: 0.6557 - val\_recall: 0.5051

Epoch 8/15

140/140 - 50s - 358ms/step - accuracy: 0.6380 - auc: 0.8123 - loss: 0.8234 - precision: 0.6827 - recall: 0.5668 - val\_accuracy: 0.6294 - val\_auc: 0.8268 - val\_loss: 0.7664 - val\_precision: 0.7037 - val\_recall: 0.5305

Epoch 9/15

140/140 - 81s - 581ms/step - accuracy: 0.6508 - auc: 0.8277 - loss: 0.7873 - precision: 0.6898 - recall: 0.5883 - val\_accuracy: 0.6510 - val\_auc: 0.8308 - val\_loss: 0.7585 - val\_precision: 0.7202 - val\_recall: 0.5520

Epoch 10/15

140/140 - 50s - 356ms/step - accuracy: 0.6584 - auc: 0.8355 - loss: 0.7691 - precision: 0.7012 - recall: 0.5992 - val\_accuracy: 0.6612 - val\_auc: 0.8425 - val\_loss: 0.7358 - val\_precision: 0.7248 - val\_recall: 0.5749

Epoch 11/15

140/140 - 49s - 351ms/step - accuracy: 0.6705 - auc: 0.8421 - loss: 0.7534 - precision: 0.7090 - recall: 0.6037 - val\_accuracy: 0.6713 - val\_auc: 0.8506 - val\_loss: 0.7209 - val\_precision: 0.7400 - val\_recall: 0.5850

Epoch 12/15

140/140 - 83s - 592ms/step - accuracy: 0.6884 - auc: 0.8522 - loss: 0.7309 - precision: 0.7240 - recall: 0.6263 - val\_accuracy: 0.6802 - val\_auc: 0.8551 - val\_loss: 0.7103 - val\_precision: 0.7425 - val\_recall: 0.5964

Epoch 13/15

140/140 - 50s - 360ms/step - accuracy: 0.6871 - auc: 0.8546 - loss: 0.7234 - precision: 0.7205 - recall: 0.6248 - val\_accuracy: 0.6865 - val\_auc: 0.8608 - val\_loss: 0.6985 - val\_precision: 0.7469 - val\_recall: 0.6104

Epoch 14/15

140/140 - 46s - 325ms/step - accuracy: 0.7065 - auc: 0.8656 - loss: 0.6961 - precision: 0.7366 - recall: 0.6514 - val\_accuracy: 0.7043 - val\_auc: 0.8699 - val\_loss: 0.6781 - val\_precision: 0.7577 - val\_recall: 0.6269

Epoch 15/15

140/140 - 86s - 617ms/step - accuracy: 0.7065 - auc: 0.8670 - loss: 0.6909 - precision: 0.7408 - recall: 0.6503 - val\_accuracy: 0.7018 - val\_auc: 0.8741 - val\_loss: 0.6695 - val\_precision: 0.7602 -

val\_recall: 0.6358

Training time for DenseNet121: 2402.36 seconds

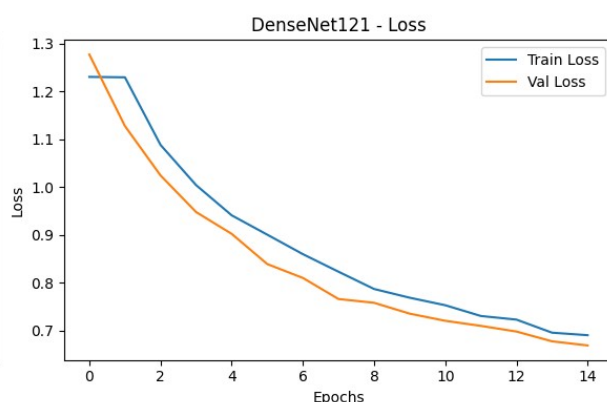


\n Evaluating DenseNet121...

1/1 \_\_\_\_\_ 12s 12s/step  
1/1 \_\_\_\_\_ 0s 139ms/step  
1/1 \_\_\_\_\_ 0s 143ms/step  
1/1 \_\_\_\_\_ 0s 134ms/step  
1/1 \_\_\_\_\_ 0s 133ms/step  
1/1 \_\_\_\_\_ 0s 136ms/step  
1/1 \_\_\_\_\_ 0s 134ms/step  
1/1 \_\_\_\_\_ 0s 135ms/step  
1/1 \_\_\_\_\_ 0s 134ms/step  
1/1 \_\_\_\_\_ 0s 133ms/step  
1/1 \_\_\_\_\_ 0s 133ms/step  
1/1 \_\_\_\_\_ 0s 131ms/step  
1/1 \_\_\_\_\_ 0s 136ms/step  
1/1 \_\_\_\_\_ 0s 156ms/step  
1/1 \_\_\_\_\_ 0s 154ms/step  
1/1 \_\_\_\_\_ 0s 154ms/step  
1/1 \_\_\_\_\_ 0s 154ms/step  
1/1 \_\_\_\_\_ 0s 154ms/step  
1/1 \_\_\_\_\_ 0s 151ms/step  
1/1 \_\_\_\_\_ 12s 12s/step

Testing time: 269.43 seconds

Classification Report:

	precision	recall	f1-score	support
WithMask	0.65	0.97	0.77	242
WithoutMask	0.94	0.27	0.42	234
MaskWearedIncorrect	0.55	0.72	0.63	148
accuracy			0.65	624
macro avg	0.71	0.65	0.61	624
weighted avg	0.73	0.65	0.61	624



```

\n Training Custom CNN...
Epoch 1/15
140/140 - 58s - 416ms/step - accuracy: 0.5253 - auc: 0.6735 - loss:
1.0133 - precision: 0.5494 - recall: 0.2218 - val_accuracy: 0.4848 -
val_auc: 0.5928 - val_loss: 1.2514 - val_precision: 0.4848 -
val_recall: 0.4848
Epoch 2/15
140/140 - 70s - 499ms/step - accuracy: 0.4850 - auc: 0.5666 - loss:
1.1485 - precision: 0.4055 - recall: 0.0961 - val_accuracy: 0.4848 -
val_auc: 0.5951 - val_loss: 1.0761 - val_precision: 0.4784 -
val_recall: 0.4505
Epoch 3/15
140/140 - 82s - 586ms/step - accuracy: 0.4848 - auc: 0.5821 - loss:
1.0810 - precision: 0.4468 - recall: 0.0903 - val_accuracy: 0.4848 -

```

val\_auc: 0.6030 - val\_loss: 1.0770 - val\_precision: 0.4837 -  
val\_recall: 0.4695  
Epoch 4/15  
140/140 - 48s - 345ms/step - accuracy: 0.4877 - auc: 0.6106 - loss:  
1.0679 - precision: 0.4910 - recall: 0.1102 - val\_accuracy: 0.4860 -  
val\_auc: 0.6236 - val\_loss: 1.0604 - val\_precision: 0.4909 -  
val\_recall: 0.4467  
Epoch 5/15  
140/140 - 48s - 339ms/step - accuracy: 0.4931 - auc: 0.6288 - loss:  
1.0513 - precision: 0.5548 - recall: 0.1485 - val\_accuracy: 0.4911 -  
val\_auc: 0.6461 - val\_loss: 1.0436 - val\_precision: 0.5251 -  
val\_recall: 0.4112  
Epoch 6/15  
140/140 - 49s - 349ms/step - accuracy: 0.5011 - auc: 0.6551 - loss:  
1.0326 - precision: 0.6097 - recall: 0.1980 - val\_accuracy: 0.4937 -  
val\_auc: 0.6582 - val\_loss: 1.0421 - val\_precision: 0.5295 -  
val\_recall: 0.4213  
Epoch 7/15  
140/140 - 43s - 307ms/step - accuracy: 0.4957 - auc: 0.6617 - loss:  
1.0271 - precision: 0.6186 - recall: 0.2343 - val\_accuracy: 0.4962 -  
val\_auc: 0.6593 - val\_loss: 1.0507 - val\_precision: 0.5263 -  
val\_recall: 0.4315  
Epoch 8/15  
140/140 - 88s - 627ms/step - accuracy: 0.4998 - auc: 0.6619 - loss:  
1.0294 - precision: 0.6153 - recall: 0.2361 - val\_accuracy: 0.4898 -  
val\_auc: 0.6619 - val\_loss: 1.0465 - val\_precision: 0.5398 -  
val\_recall: 0.4213  
Epoch 9/15  
140/140 - 77s - 548ms/step - accuracy: 0.5076 - auc: 0.6686 - loss:  
1.0248 - precision: 0.6228 - recall: 0.2534 - val\_accuracy: 0.4898 -  
val\_auc: 0.6701 - val\_loss: 1.0371 - val\_precision: 0.5453 -  
val\_recall: 0.4201  
Epoch 10/15  
140/140 - 44s - 315ms/step - accuracy: 0.5060 - auc: 0.6785 - loss:  
1.0142 - precision: 0.6296 - recall: 0.2688 - val\_accuracy: 0.4924 -  
val\_auc: 0.6742 - val\_loss: 1.0299 - val\_precision: 0.5561 -  
val\_recall: 0.4150  
Epoch 11/15  
140/140 - 49s - 348ms/step - accuracy: 0.5090 - auc: 0.6805 - loss:  
1.0114 - precision: 0.6325 - recall: 0.2968 - val\_accuracy: 0.4911 -  
val\_auc: 0.6753 - val\_loss: 1.0339 - val\_precision: 0.5517 -  
val\_recall: 0.4201  
Epoch 12/15  
140/140 - 80s - 570ms/step - accuracy: 0.5096 - auc: 0.6803 - loss:  
1.0124 - precision: 0.6239 - recall: 0.3047 - val\_accuracy: 0.4886 -  
val\_auc: 0.6767 - val\_loss: 1.0363 - val\_precision: 0.5415 -  
val\_recall: 0.4226  
Epoch 13/15  
140/140 - 44s - 311ms/step - accuracy: 0.5141 - auc: 0.6856 - loss:

1.0094 - precision: 0.6301 - recall: 0.3060 - val\_accuracy: 0.4911 -  
val\_auc: 0.6802 - val\_loss: 1.0243 - val\_precision: 0.5621 -  
val\_recall: 0.4137

Epoch 14/15

140/140 - 48s - 344ms/step - accuracy: 0.5155 - auc: 0.6869 - loss:  
1.0049 - precision: 0.6306 - recall: 0.3163 - val\_accuracy: 0.4924 -  
val\_auc: 0.6806 - val\_loss: 1.0230 - val\_precision: 0.5613 -  
val\_recall: 0.4124

Epoch 15/15

140/140 - 80s - 573ms/step - accuracy: 0.5132 - auc: 0.6868 - loss:  
1.0059 - precision: 0.6293 - recall: 0.3194 - val\_accuracy: 0.4860 -  
val\_auc: 0.6818 - val\_loss: 1.0324 - val\_precision: 0.5477 -  
val\_recall: 0.4226

Training time for Custom CNN: 907.24 seconds




\n Evaluating Custom CNN...

1/1 \_\_\_\_\_ 1s 701ms/step  
1/1 \_\_\_\_\_ 0s 91ms/step  
1/1 \_\_\_\_\_ 0s 111ms/step  
1/1 \_\_\_\_\_ 0s 94ms/step  
1/1 \_\_\_\_\_ 0s 90ms/step  
1/1 \_\_\_\_\_ 0s 123ms/step  
1/1 \_\_\_\_\_ 0s 198ms/step  
1/1 \_\_\_\_\_ 0s 202ms/step  
1/1 \_\_\_\_\_ 0s 169ms/step  
1/1 \_\_\_\_\_ 0s 161ms/step  
1/1 \_\_\_\_\_ 0s 202ms/step  
1/1 \_\_\_\_\_ 0s 128ms/step  
1/1 \_\_\_\_\_ 0s 126ms/step  
1/1 \_\_\_\_\_ 0s 124ms/step  
1/1 \_\_\_\_\_ 0s 114ms/step  
1/1 \_\_\_\_\_ 0s 97ms/step  
1/1 \_\_\_\_\_ 0s 87ms/step  
1/1 \_\_\_\_\_ 0s 75ms/step  
1/1 \_\_\_\_\_ 0s 53ms/step  
1/1 \_\_\_\_\_ 1s 759ms/step

Testing time: 8.33 seconds

Classification Report:

	precision	recall	f1-score	support
WithMask	0.39	0.93	0.55	242
WithoutMask	0.10	0.00	0.01	234
MaskWearedIncorrect	0.43	0.10	0.16	148
accuracy			0.38	624
macro avg	0.31	0.34	0.24	624
weighted avg	0.29	0.38	0.25	624



```

\n Training MobileNetV2...
Epoch 1/15
140/140 - 74s - 526ms/step - accuracy: 0.4987 - auc: 0.6950 - loss:
1.0723 - precision: 0.5254 - recall: 0.4330 - val_accuracy: 0.4975 -
val_auc: 0.6939 - val_loss: 1.1075 - val_precision: 0.5392 -
val_recall: 0.4708
Epoch 2/15
140/140 - 60s - 428ms/step - accuracy: 0.5325 - auc: 0.7203 - loss:
1.0117 - precision: 0.5793 - recall: 0.4581 - val_accuracy: 0.5317 -
val_auc: 0.7514 - val_loss: 0.9305 - val_precision: 0.6164 -
val_recall: 0.4772
Epoch 3/15
140/140 - 48s - 342ms/step - accuracy: 0.6017 - auc: 0.7855 - loss:
0.8788 - precision: 0.6504 - recall: 0.5293 - val_accuracy: 0.6015 -

```

```
val_auc: 0.7955 - val_loss: 0.8443 - val_precision: 0.6606 -  
val_recall: 0.5089  
Epoch 4/15  
140/140 - 44s - 313ms/step - accuracy: 0.6306 - auc: 0.8081 - loss:  
0.8342 - precision: 0.6667 - recall: 0.5627 - val_accuracy: 0.6624 -  
val_auc: 0.8299 - val_loss: 0.7772 - val_precision: 0.7129 -  
val_recall: 0.5546  
Epoch 5/15  
140/140 - 44s - 315ms/step - accuracy: 0.6611 - auc: 0.8367 - loss:  
0.7688 - precision: 0.6967 - recall: 0.5995 - val_accuracy: 0.6853 -  
val_auc: 0.8417 - val_loss: 0.7543 - val_precision: 0.7237 -  
val_recall: 0.5850  
Epoch 6/15  
140/140 - 48s - 345ms/step - accuracy: 0.6776 - auc: 0.8502 - loss:  
0.7335 - precision: 0.7159 - recall: 0.6243 - val_accuracy: 0.7183 -  
val_auc: 0.8687 - val_loss: 0.6912 - val_precision: 0.7640 -  
val_recall: 0.6409  
Epoch 7/15  
140/140 - 52s - 375ms/step - accuracy: 0.7050 - auc: 0.8643 - loss:  
0.6997 - precision: 0.7382 - recall: 0.6508 - val_accuracy: 0.7246 -  
val_auc: 0.8733 - val_loss: 0.6787 - val_precision: 0.7595 -  
val_recall: 0.6574  
Epoch 8/15  
140/140 - 78s - 554ms/step - accuracy: 0.7117 - auc: 0.8710 - loss:  
0.6818 - precision: 0.7420 - recall: 0.6662 - val_accuracy: 0.7221 -  
val_auc: 0.8777 - val_loss: 0.6725 - val_precision: 0.7674 -  
val_recall: 0.6701  
Epoch 9/15  
140/140 - 81s - 576ms/step - accuracy: 0.7074 - auc: 0.8718 - loss:  
0.6784 - precision: 0.7388 - recall: 0.6658 - val_accuracy: 0.7386 -  
val_auc: 0.8850 - val_loss: 0.6520 - val_precision: 0.7763 -  
val_recall: 0.6827  
Epoch 10/15  
140/140 - 84s - 599ms/step - accuracy: 0.7204 - auc: 0.8837 - loss:  
0.6444 - precision: 0.7476 - recall: 0.6756 - val_accuracy: 0.7513 -  
val_auc: 0.8897 - val_loss: 0.6364 - val_precision: 0.7779 -  
val_recall: 0.6980  
Epoch 11/15  
140/140 - 77s - 553ms/step - accuracy: 0.7296 - auc: 0.8860 - loss:  
0.6395 - precision: 0.7561 - recall: 0.6891 - val_accuracy: 0.7576 -  
val_auc: 0.8977 - val_loss: 0.6145 - val_precision: 0.7938 -  
val_recall: 0.7132  
Epoch 12/15  
140/140 - 48s - 346ms/step - accuracy: 0.7359 - auc: 0.8920 - loss:  
0.6200 - precision: 0.7605 - recall: 0.6978 - val_accuracy: 0.7500 -  
val_auc: 0.8953 - val_loss: 0.6196 - val_precision: 0.7793 -  
val_recall: 0.7081  
Epoch 13/15  
140/140 - 82s - 583ms/step - accuracy: 0.7386 - auc: 0.8934 - loss:
```



0.6162 - precision: 0.7596 - recall: 0.7036 - val\_accuracy: 0.7614 -  
val\_auc: 0.9022 - val\_loss: 0.5993 - val\_precision: 0.7852 -  
val\_recall: 0.7145

Epoch 14/15

140/140 - 45s - 321ms/step - accuracy: 0.7451 - auc: 0.9004 - loss:  
0.5946 - precision: 0.7670 - recall: 0.7115 - val\_accuracy: 0.7703 -  
val\_auc: 0.9051 - val\_loss: 0.5924 - val\_precision: 0.8020 -  
val\_recall: 0.7297

Epoch 15/15

140/140 - 49s - 350ms/step - accuracy: 0.7482 - auc: 0.9022 - loss:  
0.5878 - precision: 0.7721 - recall: 0.7126 - val\_accuracy: 0.7652 -  
val\_auc: 0.9059 - val\_loss: 0.5868 - val\_precision: 0.7862 -  
val\_recall: 0.7234

Training time for MobileNetV2: 913.83 seconds

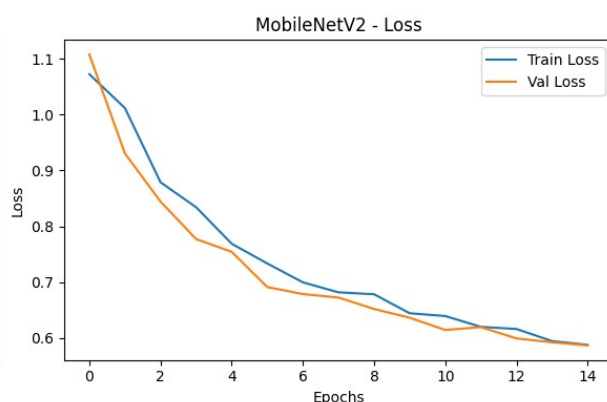


\n Evaluating MobileNetV2...

1/1 \_\_\_\_\_ 4s 4s/step  
1/1 \_\_\_\_\_ 0s 209ms/step  
1/1 \_\_\_\_\_ 0s 116ms/step  
1/1 \_\_\_\_\_ 0s 142ms/step  
1/1 \_\_\_\_\_ 0s 125ms/step  
1/1 \_\_\_\_\_ 0s 132ms/step  
1/1 \_\_\_\_\_ 0s 119ms/step  
1/1 \_\_\_\_\_ 0s 138ms/step  
1/1 \_\_\_\_\_ 0s 168ms/step  
1/1 \_\_\_\_\_ 0s 243ms/step  
1/1 \_\_\_\_\_ 0s 210ms/step  
1/1 \_\_\_\_\_ 0s 191ms/step  
1/1 \_\_\_\_\_ 0s 214ms/step  
1/1 \_\_\_\_\_ 0s 155ms/step  
1/1 \_\_\_\_\_ 0s 161ms/step  
1/1 \_\_\_\_\_ 0s 133ms/step  
1/1 \_\_\_\_\_ 0s 159ms/step  
1/1 \_\_\_\_\_ 0s 138ms/step  
1/1 \_\_\_\_\_ 0s 71ms/step  
1/1 \_\_\_\_\_ 3s 3s/step

Testing time: 16.46 seconds

Classification Report:

	precision	recall	f1-score	support
WithMask	0.64	0.98	0.78	242
WithoutMask	0.97	0.39	0.55	234
MaskWearIncorrect	0.61	0.65	0.63	148
accuracy			0.68	624
macro avg	0.74	0.67	0.65	624
weighted avg	0.76	0.68	0.66	624



```

\n Training VGG16...
Epoch 1/15
140/140 - 86s - 613ms/step - accuracy: 0.4328 - auc: 0.6011 - loss:
1.0734 - precision: 0.5085 - recall: 0.1333 - val_accuracy: 0.4848 -
val_auc: 0.6287 - val_loss: 1.0510 - val_precision: 0.4841 -
val_recall: 0.3680
Epoch 2/15
140/140 - 53s - 381ms/step - accuracy: 0.4628 - auc: 0.6138 - loss:
1.0694 - precision: 0.4911 - recall: 0.2092 - val_accuracy: 0.4848 -
val_auc: 0.6588 - val_loss: 1.0409 - val_precision: 0.4895 -
val_recall: 0.4416
Epoch 3/15
140/140 - 53s - 381ms/step - accuracy: 0.4691 - auc: 0.6178 - loss:
1.0671 - precision: 0.4702 - recall: 0.1960 - val_accuracy: 0.4848 -

```

val\_auc: 0.6837 - val\_loss: 1.0278 - val\_precision: 0.4966 -  
val\_recall: 0.4670  
Epoch 4/15  
140/140 - 53s - 376ms/step - accuracy: 0.4727 - auc: 0.6373 - loss:  
1.0505 - precision: 0.4948 - recall: 0.2126 - val\_accuracy: 0.4848 -  
val\_auc: 0.7018 - val\_loss: 1.0140 - val\_precision: 0.5155 -  
val\_recall: 0.4645  
Epoch 5/15  
140/140 - 48s - 343ms/step - accuracy: 0.4973 - auc: 0.6602 - loss:  
1.0286 - precision: 0.5447 - recall: 0.2348 - val\_accuracy: 0.4848 -  
val\_auc: 0.7136 - val\_loss: 1.0018 - val\_precision: 0.5396 -  
val\_recall: 0.4492  
Epoch 6/15  
140/140 - 85s - 610ms/step - accuracy: 0.4973 - auc: 0.6736 - loss:  
1.0160 - precision: 0.5624 - recall: 0.2493 - val\_accuracy: 0.4848 -  
val\_auc: 0.7279 - val\_loss: 0.9881 - val\_precision: 0.5800 -  
val\_recall: 0.4416  
Epoch 7/15  
140/140 - 53s - 382ms/step - accuracy: 0.5038 - auc: 0.6869 - loss:  
1.0040 - precision: 0.5870 - recall: 0.2540 - val\_accuracy: 0.4848 -  
val\_auc: 0.7385 - val\_loss: 0.9765 - val\_precision: 0.6049 -  
val\_recall: 0.4353  
Epoch 8/15  
140/140 - 53s - 381ms/step - accuracy: 0.5137 - auc: 0.6964 - loss:  
0.9942 - precision: 0.5975 - recall: 0.2760 - val\_accuracy: 0.4848 -  
val\_auc: 0.7496 - val\_loss: 0.9653 - val\_precision: 0.6225 -  
val\_recall: 0.4289  
Epoch 9/15  
140/140 - 51s - 361ms/step - accuracy: 0.5206 - auc: 0.7132 - loss:  
0.9787 - precision: 0.6258 - recall: 0.2776 - val\_accuracy: 0.4848 -  
val\_auc: 0.7601 - val\_loss: 0.9538 - val\_precision: 0.6523 -  
val\_recall: 0.4213  
Epoch 10/15  
140/140 - 81s - 582ms/step - accuracy: 0.5300 - auc: 0.7260 - loss:  
0.9667 - precision: 0.6357 - recall: 0.2928 - val\_accuracy: 0.4848 -  
val\_auc: 0.7729 - val\_loss: 0.9431 - val\_precision: 0.6543 -  
val\_recall: 0.4251  
Epoch 11/15  
140/140 - 53s - 377ms/step - accuracy: 0.5459 - auc: 0.7371 - loss:  
0.9547 - precision: 0.6476 - recall: 0.2988 - val\_accuracy: 0.4860 -  
val\_auc: 0.7743 - val\_loss: 0.9333 - val\_precision: 0.6800 -  
val\_recall: 0.4099  
Epoch 12/15  
140/140 - 49s - 351ms/step - accuracy: 0.5511 - auc: 0.7445 - loss:  
0.9467 - precision: 0.6679 - recall: 0.3109 - val\_accuracy: 0.4924 -  
val\_auc: 0.7853 - val\_loss: 0.9230 - val\_precision: 0.6792 -  
val\_recall: 0.4137  
Epoch 13/15  
140/140 - 84s - 598ms/step - accuracy: 0.5578 - auc: 0.7566 - loss:

0.9345 - precision: 0.6824 - recall: 0.3239 - val\_accuracy: 0.4962 -  
val\_auc: 0.7884 - val\_loss: 0.9139 - val\_precision: 0.6799 -  
val\_recall: 0.4124

Epoch 14/15

140/140 - 54s - 383ms/step - accuracy: 0.5681 - auc: 0.7628 - loss:  
0.9252 - precision: 0.6763 - recall: 0.3286 - val\_accuracy: 0.5013 -  
val\_auc: 0.7934 - val\_loss: 0.9057 - val\_precision: 0.6728 -  
val\_recall: 0.4150

Epoch 15/15

140/140 - 54s - 383ms/step - accuracy: 0.5737 - auc: 0.7691 - loss:  
0.9169 - precision: 0.6922 - recall: 0.3562 - val\_accuracy: 0.5127 -  
val\_auc: 0.7949 - val\_loss: 0.8969 - val\_precision: 0.7007 -  
val\_recall: 0.4099

Training time for VGG16: 910.18 seconds

\n Evaluating VGG16...

1/1 \_\_\_\_\_ 1s 1s/step  
1/1 \_\_\_\_\_ 0s 212ms/step  
1/1 \_\_\_\_\_ 0s 264ms/step  
1/1 \_\_\_\_\_ 0s 240ms/step  
1/1 \_\_\_\_\_ 0s 210ms/step  
1/1 \_\_\_\_\_ 0s 194ms/step  
1/1 \_\_\_\_\_ 0s 194ms/step  
1/1 \_\_\_\_\_ 0s 189ms/step  
1/1 \_\_\_\_\_ 0s 199ms/step  
1/1 \_\_\_\_\_ 0s 190ms/step  
1/1 \_\_\_\_\_ 0s 209ms/step  
1/1 \_\_\_\_\_ 0s 213ms/step  
1/1 \_\_\_\_\_ 0s 203ms/step  
1/1 \_\_\_\_\_ 0s 219ms/step  
1/1 \_\_\_\_\_ 0s 275ms/step  
1/1 \_\_\_\_\_ 0s 221ms/step  
1/1 \_\_\_\_\_ 0s 210ms/step  
1/1 \_\_\_\_\_ 0s 203ms/step  
1/1 \_\_\_\_\_ 0s 202ms/step  
1/1 \_\_\_\_\_ 1s 1s/step

Testing time: 9.80 seconds

Classification Report:

	precision	recall	f1-score	support
WithMask	0.39	1.00	0.56	242
WithoutMask	1.00	0.01	0.02	234
MaskWearredIncorrect	0.00	0.00	0.00	148
accuracy			0.39	624
macro avg	0.46	0.34	0.19	624
weighted avg	0.53	0.39	0.22	624

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/  
\_classification.py:1565: UndefinedMetricWarning: Precision is ill-

```

defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))

```

